

OSM LAB - DYNAMIC PROGRAMMING

Jason DeBacker
June 18-22, 2018
Notes for Lecture #3

Today

- Discrete Choice Dynamic Programming
- Tips for better computing performance

Discrete Choice Cake Eating Problem: (an example of an optimal stopping problem)

- control: {eat cake, leave cake} \rightarrow binary (0,1 choice)
 - $z \in \{1, 0\}$
- state: $w, \varepsilon \Rightarrow$ know w and ε at the time of the decision
- transition: $w' = \rho w$ if $z = 0$ (grow/shrink leftover cake), $w' = 0$ if $z = 1$ (cake eaten in period 1, no w')
- value function: $V(w, \varepsilon) = \max\{\underbrace{V^0(w, \varepsilon)}_{\text{leave cake}}, \underbrace{V^1(w, \varepsilon)}_{\text{eat cake}}\}, \forall (w, \varepsilon)$
 - $V^0(w, \varepsilon) = \beta E_{\varepsilon'|\varepsilon} V(\rho w, \varepsilon')$
 - $V^1(w, \varepsilon) = \varepsilon u(w)$
- policy function: $z(w, \varepsilon) \in \{0, 1\}, \forall (w, \varepsilon)$
- Choice depends on:
 - State variables: (w & ε) b/c in state vector
 - Parameters:
 - * ρ , b/c as $\rho \uparrow$, gain to waiting
 - * β , $\beta \downarrow$ cost to waiting
 - * Π : the transition matrix
- NOTE: No Euler equation in discrete case - eat or don't eat - it's not continuous
- e.g., $\rho = 1, \varepsilon \in \{\varepsilon_L, \varepsilon_H\}$
 - $z(w, \varepsilon_H) = 1, \forall w$: nothing to wait for!
 - $z(w, \varepsilon_L) = \{0, 1\} \rightarrow$ wait if: β near 1 or π_{LH} sufficiently high
 - * NOTE: w unimportant b/c its in both V^0 and V^1 decisions
 - * How high does π_{LH} have to be to wait?
 - * wait if: $\varepsilon_L u(w) \leq \beta \{E_{\varepsilon'|\varepsilon_L} V(w, \varepsilon')\} = \beta \{\pi_{LH} \varepsilon_H u(w) + \pi_{LL} V(w, \varepsilon_L)\}$
 - * B/c always eat in high, and assuming never eat in low (this is the RHS of the equality), know that: $E_{\varepsilon'|\varepsilon_L} V(w, \varepsilon') = \pi_{LH} \varepsilon_H u(w) + \pi_{LL} \beta E_{\varepsilon'|\varepsilon_L} V(w, \varepsilon')$
 - * Solving for $V(w, \varepsilon') \Rightarrow E_{\varepsilon'|\varepsilon_L} V(w, \varepsilon') = \frac{\pi_{LH} \varepsilon_H u(w)}{1 - \beta \pi_{LL}}$
 - * Thus, wait if $\varepsilon_L u(w) \leq \beta \{E_{\varepsilon'|\varepsilon_L} V(w, \varepsilon')\} = \frac{\beta \pi_{LH} \varepsilon_H u(w)}{1 - \beta \pi_{LL}}$
 - * Note that we can divide both sides by $u(w)$: $\varepsilon_L \leq \frac{\beta \pi_{LH} \varepsilon_H}{1 - \beta \pi_{LL}}$
 - * So, without growth in the size of the cake overtime, the decision rule is not a function of the size of the cake or the parameterization of the utility function.

* NOTE: this is not the case if the size of the cake is growing.

Some final tips on speeding up computations

- In Python, use `Numba` and its just-in-time compilation to speed up loops
 - Otherwise, in Python or other higher-level languages, “vectorize” your code so that you are doing operations on arrays rather than through loops with element-by-element operations
- Consider parallel processing (e.g., to evaluate functions are different parts of your state space simultaneously)
 - The `dask` package for Python makes multiprocessing relatively straightforward